

```
In [4]: import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
```

```
In [5]: #Load the dataset
df = pd.read_csv("IMDB Top 250 Movies.csv")
df.head()
```

```
Out[5]:
```

| | rank | name | year | rating | genre | certificate | run_time | tagline | |
|---|------|--------------------------|------|--------|--------------------|-------------|----------|---|----|
| 0 | 1 | The Shawshank Redemption | 1994 | 9.3 | Drama | R | 2h 22m | Fear can hold you prisoner. Hope can set you f... | 2 |
| 1 | 2 | The Godfather | 1972 | 9.2 | Crime,Drama | R | 2h 55m | An offer you can't refuse. | |
| 2 | 3 | The Dark Knight | 2008 | 9.0 | Action,Crime,Drama | PG-13 | 2h 32m | Why So Serious? | 18 |
| 3 | 4 | The Godfather Part II | 1974 | 9.0 | Crime,Drama | R | 3h 22m | All the power on earth can't change destiny. | 1 |
| 4 | 5 | 12 Angry Men | 1957 | 9.0 | Crime,Drama | Approved | 1h 36m | Life Is In Their Hands - Death Is On Their Mi... | |

Getting basic info about the data

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 250 entries, 0 to 249
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   rank                   250 non-null   int64
1   name                   250 non-null   object
2   year                   250 non-null   int64
3   rating                 250 non-null   float64
4   genre                  250 non-null   object
5   certificate            250 non-null   object
6   run_time               250 non-null   object
7   tagline                250 non-null   object
8   budget                 250 non-null   object
9   box_office             250 non-null   object
10  casts                  250 non-null   object
11  directors              250 non-null   object
12  writers                250 non-null   object
dtypes: float64(1), int64(2), object(10)
memory usage: 25.5+ KB
```

```
In [7]: df.describe()
```

```
Out[7]:
```

| | rank | year | rating |
|--------------|------------|-------------|------------|
| count | 250.000000 | 250.000000 | 250.000000 |
| mean | 125.500000 | 1986.360000 | 8.307200 |
| std | 72.312977 | 25.125356 | 0.229081 |
| min | 1.000000 | 1921.000000 | 8.000000 |
| 25% | 63.250000 | 1966.250000 | 8.100000 |
| 50% | 125.500000 | 1994.000000 | 8.200000 |
| 75% | 187.750000 | 2006.000000 | 8.400000 |
| max | 250.000000 | 2022.000000 | 9.300000 |

Feature Engineering

Budget

```
In [8]: set(df['budget'])
```

```
Out[8]: {'$3300000',
'$8240000',
'1000000',
'10000000',
'100000000',
'102000000',
'103000000',
}
```

'11000000',
'11400000',
'11500000',
'115000000',
'1200000',
'12000000',
'125000000',
'1288000',
'1300000',
'13000000',
'133000',
'13500000',
'1400000',
'140000000',
'14400000',
'14600000',
'150000',
'1500000',
'15000000',
'150000000',
'160000000',
'16400000',
'165000000',
'170000000',
'17500000',
'175000000',
'1752000',
'1800000',
'18000000',
'180000000',
'18500000',
'185000000',
'1900000',
'19000000',
'2000000',
'20000000',
'200000000',
'2100000',
'2200000',
'22000000',
'229575',
'23000000',
'24000000',
'2400000000',
'2479000',
'250000',
'2500000',
'25000000',
'250000000',
'2540800',
'2777000',
'2800000',
'28000000',
'2883848',
'3000000',
'30000000',
'3101000',

'31500000',
'3180000',
'3200000',
'321000000',
'325000',
'32500000',
'3300000',
'33000000',
'350000',
'35000000',
'356000000',
'3700000',
'3800000',
'38000000',
'3977000',
'4000000',
'40000000',
'4500000',
'45000000',
'46000000',
'48000000',
'500000',
'5000000',
'52000000',
'54000000',
'5500000',
'55000000',
'550000000',
'58000000',
'600000',
'6000000',
'60000000',
'63000000',
'6400000',
'6800000',
'7000000',
'70000000',
'72000000',
'750000',
'800000',
'8000000',
'80000000',
'806947',
'8200000',
'839727',
'9000000',
'90000000',
'910000',
'92000000',
'923000',
'927262',
'93000000',
'935000',
'94000000',
'950000',
'9500000',
'95000000',

```
'960000',
'97000000',
'97600000',
'EM\xa032000000',
'EM\xa06000000',
'Not Available',
'RF\xa0115000000'}
```

In [9]: *#converting all the currency in dollar and only keeping numeric value*

```
def convert_currency(value):
    if "DEM" in value:
        value = value.replace('DEM\xa0', '')
        value = float(value)*0.54
    elif "FRF" in value:
        value = value.replace('FRF\xa0', '')
        value = float(value)*0.16
    return value
```

In [10]: *#removing the \$ sign from the value*

```
df['budget'] = df['budget'].str.replace('$', '')
df['budget'] = df['budget'].apply(convert_currency)
df['budget'] = pd.to_numeric(df['budget'],
                             errors='coerce',
                             downcast='float')
```

/tmp/ipykernel_138/4123476812.py:3: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expressions will *not* be treated as literal strings when regex=True.

```
df['budget'] = df['budget'].str.replace('$', '')
```

Box Office Collection

In [11]: `set(df['box_office'])`

```
Out[11]: {'1006234167',
'101209702',
'104880868',
'106285522',
'1067316101',
'1074458282',
'1081169825',
'109114817',
'109676311',
'1098',
'1109802321',
'1146457748',
'11477',
'11487676',
'116112375',
'1166717',
'117250402',
'120072577',
```

'12180',
'122126687',
'1226507',
'126216940',
'127244',
'13019063',
'131060248',
'1342359942',
'134686457',
'1349711',
'135342',
'138545632',
'141603197',
'142319',
'144738',
'14480',
'14651',
'1488732821',
'15000000',
'151052',
'15421226',
'156000000',
'159167799',
'159428329',
'162861289',
'1645133',
'16767475',
'170005875',
'171627434',
'173082189',
'1740429',
'174122191',
'180563636',
'180906076',
'187436818',
'187733202',
'18778738',
'1921847111',
'1940906',
'195088',
'19569225',
'19632715',
'199632',
'2000288',
'2052415039',
'20745728',
'20908467',
'213216216',
'213928762',
'216639112',
'216763646',
'21897373',
'225508210',
'228178',
'22926076',
'230098753',
'23308615',

'23335817',
'23341568',
'23402427',
'235860116',
'23661347',
'23690757',
'237536126',
'23875127',
'24427162',
'250341816',
'25253887',
'25637669',
'26023860',
'262676096',
'264118201',
'26960374',
'269958228',
'27200000',
'272742922',
'2799439100',
'28570902',
'286801374',
'28884504',
'2913644',
'291480452',
'29225935',
'29328',
'294805697',
'300073',
'30680793',
'311212',
'31207',
'316791257',
'32052925',
'321457747',
'321752656',
'322161245',
'327333559',
'33882243',
'346258',
'35401758',
'35566',
'355822319',
'356296601',
'37034579',
'373672993',
'3753929',
'383336762',
'384479940',
'389925971',
'394436586',
'40047078',
'402382193',
'406878233',
'415261382',
'4164283',
'41722424',

'41960',
'424208848',
'426074373',
'426588510',
'436655',
'441306145',
'45720631',
'4626532',
'46358827',
'463618',
'467222728',
'46808',
'46969409',
'47036784',
'47335804',
'474171806',
'475106177',
'476512065',
'47961919',
'482349603',
'49074379',
'49396747',
'494879471',
'5014000',
'503162313',
'50419',
'516962',
'52066791',
'520881154',
'521311890',
'52287414',
'5252',
'52767889',
'538375067',
'5435024',
'5473337',
'56675895',
'579707738',
'5987386',
'599146',
'60262836',
'60418',
'60611975',
'619179950',
'623726085',
'631607053',
'654264015',
'65884703',
'65889846',
'67',
'672806432',
'678226465',
'6788659',
'71108591',
'72275',
'735099102',
'7390108',


```
'74036715',
'74437720',
'750000 (estimated)',
'76182388',
'7693',
'77356942',
'773867216',
'775398007',
'7798146',
'78371200',
'800000 (estimated)',
'81379',
'814337054',
'83557872',
'836848102',
'83862032',
'8574081',
'858848019',
'898204420',
'910000 (estimated)',
'92181574',
'941637960',
'947944270',
'9523464',
'955',
'962002',
'96302',
'968511805',
'96983009',
'970263',
'98690254',
'Not Available'}
```

In [12]: *#removing the estimated and taking float value*

```
df['box_office'] = df['box_office'].str.replace('(estimated)', '')
df['box_office'] = pd.to_numeric(df['box_office'],
                                errors='coerce',
                                downcast='integer')
```

/tmp/ipykernel_138/744310834.py:3: FutureWarning: The default value of re
gex will change from True to False in a future version.

```
df['box_office'] = df['box_office'].str.replace('(estimated)', '')
```

Run Time

In [13]: *#converting the run-time of movies into minutes*

```
def convert_to_minutes(time):  
    try:  
        total_minutes = 0  
        if 'h' in time:  
            time = time.split('h')  
            hours = int(time[0])  
            mins = int(time[1][:-1]) if time[1][:-1].strip().isdigit() else 0  
            total_minutes = hours*60 + mins  
        else:  
            total_minutes = int(time[:-1])  
        return total_minutes  
    except ValueError:  
        return float('NaN')
```

In [14]: `df['run_time'] = df['run_time'].apply(convert_to_minutes)`

In [15]: *#splitting genre, directors, writers and casts and made a list*

```
df['genre'] = df['genre'].str.split(",")    #genre  
df['directors'] = df['directors'].str.split(",")    #directors  
df['writers'] = df['writers'].str.split(",")    #writers  
df['casts'] = df['casts'].str.split(",")    #casts
```

In [16]: `df.head()`

| Out[16]: | rank | name | year | rating | genre | certificate | run_time | tagline | budget |
|----------|------|--------------------------|------|--------|------------------------|-------------|----------|---|-------------|
| 0 | 1 | The Shawshank Redemption | 1994 | 9.3 | [Drama] | R | 142.0 | Fear can hold you prisoner. Hope can set you f... | 25000000.0 |
| 1 | 2 | The Godfather | 1972 | 9.2 | [Crime, Drama] | R | 175.0 | An offer you can't refuse. | 6000000.0 |
| 2 | 3 | The Dark Knight | 2008 | 9.0 | [Action, Crime, Drama] | PG-13 | 152.0 | Why So Serious? | 185000000.0 |
| 3 | 4 | The Godfather Part II | 1974 | 9.0 | [Crime, Drama] | R | 202.0 | All the power on earth can't change destiny. | 13000000.0 |
| 4 | 5 | 12 Angry Men | 1957 | 9.0 | [Crime, Drama] | Approved | 96.0 | Life Is In Their Hands - Death Is On Their Mi... | 350000.0 |

In [17]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 250 entries, 0 to 249
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   rank            250 non-null    int64
1   name            250 non-null    object
2   year            250 non-null    int64
3   rating          250 non-null    float64
4   genre           250 non-null    object
5   certificate      250 non-null    object
6   run_time        249 non-null    float64
7   tagline         250 non-null    object
8   budget          208 non-null    float32
9   box_office      217 non-null    float64
10  casts           250 non-null    object
11  directors        250 non-null    object
12  writers         250 non-null    object
dtypes: float32(1), float64(3), int64(2), object(7)
memory usage: 24.5+ KB
```

```
In [18]: df.describe()
```

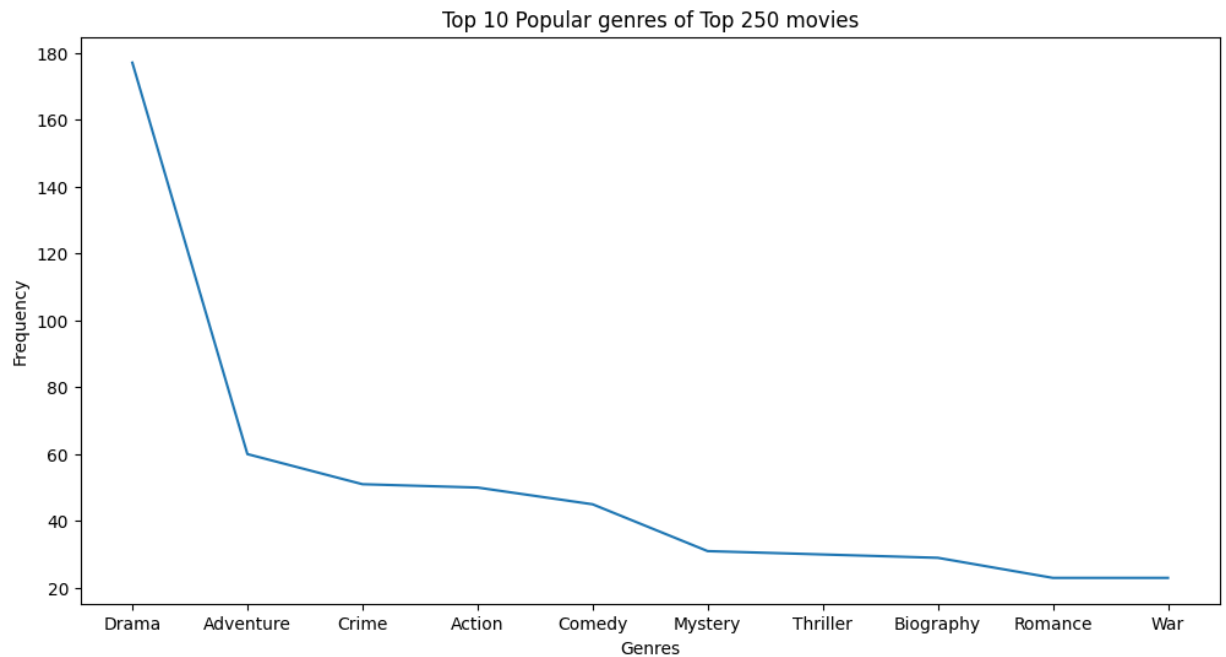
```
Out[18]:
```

| | rank | year | rating | run_time | budget | box_office |
|--------------|------------|-------------|------------|------------|--------------|--------------|
| count | 250.000000 | 250.000000 | 250.000000 | 249.000000 | 2.080000e+02 | 2.170000e+02 |
| mean | 125.500000 | 1986.360000 | 8.307200 | 129.052209 | 5.245898e+07 | 2.382076e+08 |
| std | 72.312977 | 25.125356 | 0.229081 | 30.223383 | 1.770686e+08 | 3.825392e+08 |
| min | 1.000000 | 1921.000000 | 8.000000 | 45.000000 | 1.330000e+05 | 6.700000e+01 |
| 25% | 63.250000 | 1966.250000 | 8.100000 | 107.000000 | 3.000000e+06 | 8.574081e+06 |
| 50% | 125.500000 | 1994.000000 | 8.200000 | 126.000000 | 1.500000e+07 | 7.403672e+07 |
| 75% | 187.750000 | 2006.000000 | 8.400000 | 145.000000 | 4.900000e+07 | 3.217527e+08 |
| max | 250.000000 | 2022.000000 | 9.300000 | 238.000000 | 2.400000e+09 | 2.799439e+09 |

EDA (Exploratory Data Analysis)


What are the popular genres of Top 250 movies?


```
In [19]: diff_genres_count = df['genre'].explode().value_counts()
plt.figure(figsize=(12, 6))
plt.plot(diff_genres_count[:10])
plt.xlabel("Genres")
plt.ylabel("Frequency")
plt.title("Top 10 Popular genres of Top 250 movies")
plt.show()
```



How does ratings vary across the genres?

```
In [20]: rating_per_genre = df[['genre', 'rating']].explode('genre').groupby('genre')
genre_rating_dict = rating_per_genre['rating'].apply(list).to_dict()
grouped_data = list(genre_rating_dict.values())
mean_rating_per_genre = rating_per_genre.mean()
plt.figure(figsize=(24, 6))
plt.boxplot(grouped_data)
plt.plot(range(1, len(genre_rating_dict) + 1), mean_rating_per_genre, 'o')
plt.xticks(range(1, len(genre_rating_dict) + 1), list(genre_rating_dict.keys()))
plt.xlabel("Genre")
plt.ylabel("Ratings")
plt.title("Ratings vs Genre")
plt.show()
```





What were the categories of Top 250 Movies?

There are several box office categories but there is no specific formula or threshold for determining the categories. Instead, the categories are based on industry standards and expectations, as well as the production budget and marketing costs of the movie. For the sake of simplicity, we're using the following formula in comparison to their budget.

- **All time Blockbuster:** Above 300%
- **Blockbuster:** 200%-300%
- **Super Hit:** 175%-200%
- **Hit:** 125%-175%
- **Average:** 100%-125%
- **Flop:** 75%-100%
- **Super Flop:** 50%-75%
- **Disaster:** Below 50%

```
In [21]: #categorising on the basis of ROI
#applying conditions to roi
def convert_to_categories(roi):
    if roi > 3.0:
        return "All Time Blockbuster"
    elif roi > 2.0:
        return "Blockbuster"
    elif roi > 1.75:
        return "Super Hit"
    elif roi > 1.25:
        return "Hit"
    elif roi > 1.0:
        return "Average"
    elif roi > 0.75:
        return "Flop"
    elif roi > 0.5:
        return "Super Flop"
    else:
        return "Disaster"
```

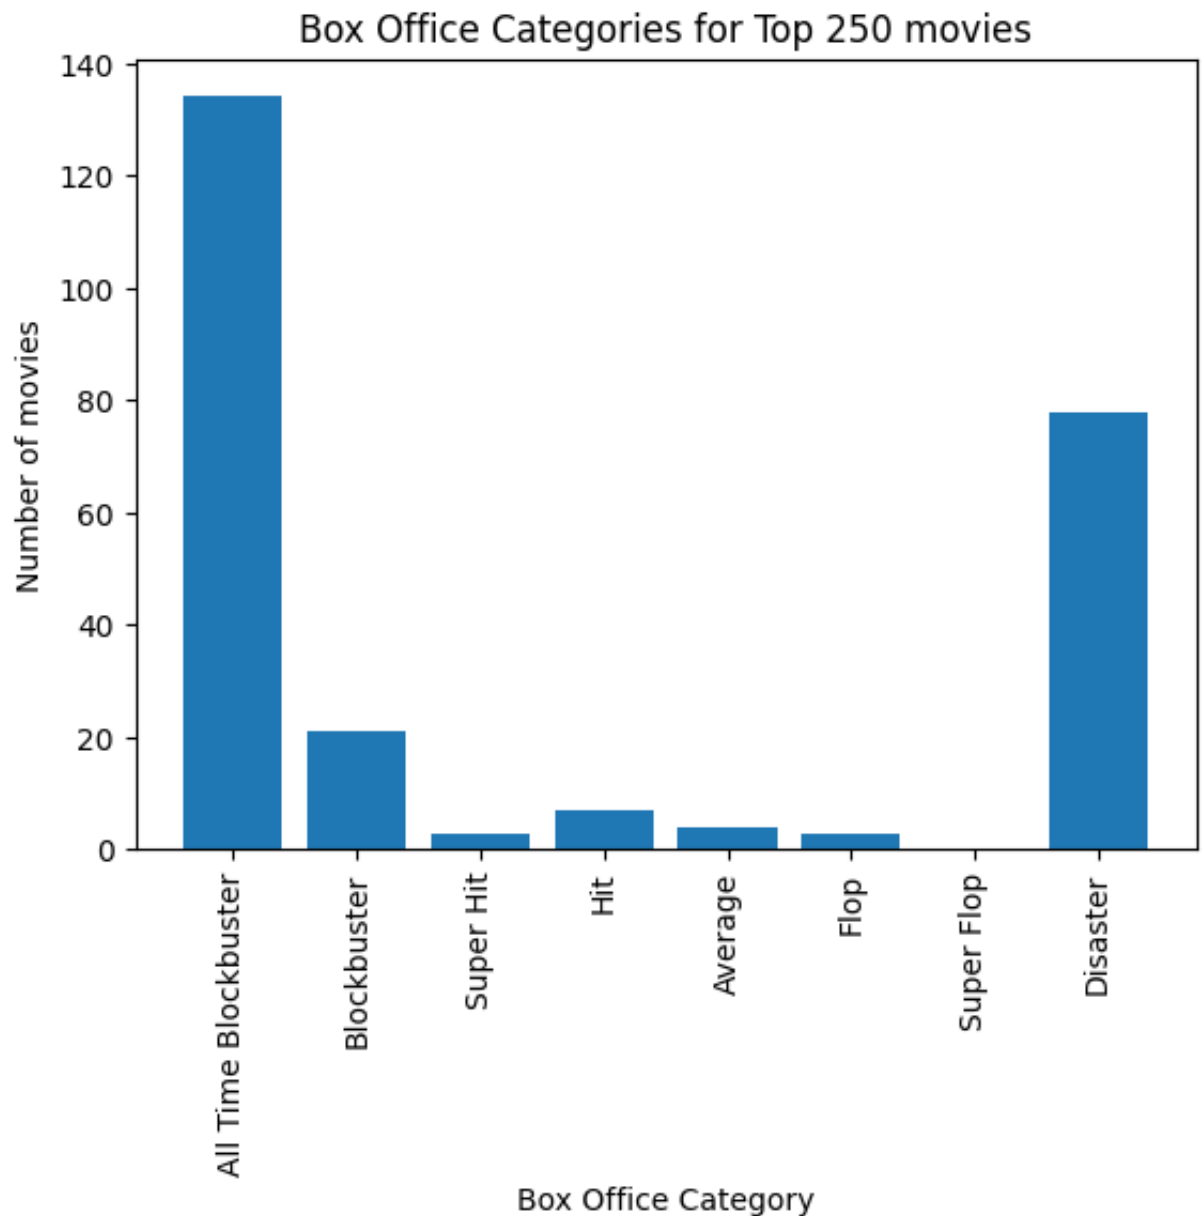
```
In [22]: #calculated roi(return of investment)
df['roi'] = df['box_office']/df['budget'] #formula for ROI
df['box-office-status'] = df['roi'].apply(convert_to_categories)
df.head()
```

Out[22]:

| | rank | name | year | rating | genre | certificate | run_time | tagline | budget |
|---|------|--------------------------|------|--------|------------------------|-------------|----------|---|-------------|
| 0 | 1 | The Shawshank Redemption | 1994 | 9.3 | [Drama] | R | 142.0 | Fear can hold you prisoner. Hope can set you f... | 25000000.0 |
| 1 | 2 | The Godfather | 1972 | 9.2 | [Crime, Drama] | R | 175.0 | An offer you can't refuse. | 6000000.0 |
| 2 | 3 | The Dark Knight | 2008 | 9.0 | [Action, Crime, Drama] | PG-13 | 152.0 | Why So Serious? | 185000000.0 |
| 3 | 4 | The Godfather Part II | 1974 | 9.0 | [Crime, Drama] | R | 202.0 | All the power on earth can't change destiny. | 13000000.0 |
| 4 | 5 | 12 Angry Men | 1957 | 9.0 | [Crime, Drama] | Approved | 96.0 | Life Is In Their Hands - Death Is On Their Mi... | 350000.0 |

```
In [23]: #plotting according to its success rate


xticks = ["All Time Blockbuster", "Blockbuster", "Super Hit", "Hit", "Ave
box_office_status = df['box-office-status'].value_counts()
box_office_status = box_office_status.reindex(xticks)
plt.bar(box_office_status.keys(), box_office_status.values)
plt.xlabel("Box Office Category")
plt.ylabel("Number of movies")
plt.xticks(rotation="vertical")
plt.title("Box Office Categories for Top 250 movies")
plt.show()
```




```
In [24]: #distribution according to the certificates


movie_certs = df['certificate'].value_counts()
others_count = movie_certs[movie_certs < 5].sum()
movie_certs = movie_certs.drop(movie_certs[movie_certs < 5].index.tolist())
movie_certs['Others'] = others_count
plt.figure(figsize=(8, 8))
plt.pie(movie_certs, \
        labels=movie_certs.keys(), \
        startangle = 90, \
        autopct='%1.1f%%', \
        counterclock=False)
plt.title("Distribution of Top 250 movies according to the certificates")
plt.show()
```

Distribution of Top 250 movies according to the certificates



In [25]: *#top-10 directors plotted*

```
best_directors = df['directors'].explode().value_counts()
plt.barh(best_directors.keys()[:10], best_directors.values[:10])
plt.ylabel("DIrectors")
plt.xlabel("Total number of movies")
plt.title("Top 10 directors of top 250 movies")
# plt.plot(best_directors[:5])
plt.show()
```



```
In [26]: dir_probs = best_directors/250
cast_probs = df['casts'].explode().value_counts()/250
writer_probs = df['writers'].explode().value_counts()/250
certs_probs = movie_certs/250
genre_probs = diff_genres_count/250

#calculated probability ... applied regression

#for genres
def convert_genres(genres):
    genre_factor = -0.06
    genre_prob = 0
    for genre in genres:
        genre_prob += genre_probs[genre] if genre in genre_probs else genre_f

    return genre_prob

#for certificates
def convert_certificate(certificate):
    return certs_probs[certificate] if certificate in certs_probs\
    else certs_probs['Others']
```

```

#for directors

def convert_directors(directors):
    dir_factor = -0.02
    dir_prob = 0
    for director in directors:
        dir_prob += dir_probs[director] if director in dir_probs else dir_factor

    return dir_prob

#for casts

def convert_casts(casts):
    cast_factor = -0.015
    cast_prob = 0
    for cast in casts:
        cast_prob += cast_probs[cast] if cast in cast_probs else cast_factor

    return cast_prob

#for writers

def convert_writers(writers):
    writer_factor = -0.007
    writer_prob = 0
    for writer in writers:
        writer_prob += writer_probs[writer] if writer in writer_probs else writer_factor

    return writer_prob

regression_df = df.copy()
regression_df['genre'] = regression_df['genre'].apply(convert_genres)
regression_df['certificate'] = regression_df['certificate'].apply(convert_certificate)
regression_df['directors'] = regression_df['directors'].apply(convert_directors)
regression_df['casts'] = regression_df['casts'].apply(convert_casts)
regression_df['writers'] = regression_df['writers'].apply(convert_writers)

regression_df['run_time'] = regression_df['run_time'].fillna(regression_df['run_time'].mean())
regression_df['budget'] = regression_df['budget'].fillna(regression_df['budget'].mean())
regression_df['box_office'] = regression_df['box_office'].fillna(regression_df['box_office'].mean())
regression_df['roi'] = regression_df['box_office']/regression_df['budget']
regression_df = regression_df.drop(columns=['name', 'rank', 'tagline', 'budget', 'box_office'])

regression_df.head()

```

Out[26]:

| | year | rating | genre | certificate | run_time | budget | box_office | casts | directors |
|---|------|--------|-------|-------------|----------|-------------|--------------|-------|-----------|
| 0 | 1994 | 9.3 | 0.708 | 0.388 | 142.0 | 25000000.0 | 2.888450e+07 | 0.116 | 0.00 |
| 1 | 1972 | 9.2 | 0.912 | 0.388 | 175.0 | 6000000.0 | 2.503418e+08 | 0.144 | 0.00 |
| 2 | 2008 | 9.0 | 1.112 | 0.140 | 152.0 | 185000000.0 | 1.006234e+09 | 0.164 | 0.00 |
| 3 | 1974 | 9.0 | 0.912 | 0.388 | 202.0 | 13000000.0 | 4.796192e+07 | 0.172 | 0.00 |
| 4 | 1957 | 9.0 | 0.912 | 0.056 | 96.0 | 350000.0 | 9.550000e+02 | 0.112 | 0.00 |

In [27]: `regression_df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 250 entries, 0 to 249
Data columns (total 11 columns):
#   Column          Non-Null Count  Dtype
---  -
0   year            250 non-null    int64
1   rating          250 non-null    float64
2   genre           250 non-null    float64
3   certificate      250 non-null    float64
4   run_time        250 non-null    float64
5   budget          250 non-null    float32
6   box_office      250 non-null    float64
7   casts           250 non-null    float64
8   directors       250 non-null    float64
9   writers         250 non-null    float64
10  roi             250 non-null    float64
dtypes: float32(1), float64(9), int64(1)
memory usage: 20.6 KB
```

In [28]: *#calculated the accuracy*

```
def count_accuracy(y_true, y_pred):
    threshold = 0.2
    count = 0
    for i in range(len(y_true)):
        if abs(y_true[i]-y_pred[i]) < threshold:
            count += 1
    return count/len(y_true)
```

Test

In [29]: *#linear regression*

```
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
#Define the model
model = LinearRegression()
y = np.array(regression_df['rating'])
x = np.array(regression_df.loc[:, regression_df.columns!='rating'])
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.15,
reg = model.fit(X_train, y_train)
y_pred = reg.predict(X_test)
accuracy = count_accuracy(y_test, y_pred)
#Printed the accuracy
print(accuracy)
```

0.6578947368421053

```
In [30]: #Support Vector Regression (SVR)

from sklearn.svm import SVR
#Define the model
model = SVR()
y = np.array(regression_df['rating'])
x = np.array(regression_df.loc[:, regression_df.columns!='rating'])
X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.15,
reg = model.fit(X_train, y_train)
y_pred = reg.predict(X_test)
accuracy = count_accuracy(y_test, y_pred)
#Printed the accuracy
print(accuracy)

0.7631578947368421
```

In []: